

Claude Code for Accounting Research

Module 6 — Seminar: AI, the Production of Knowledge, and
the Academic Research Market

John Barrios


Yale School of Management

The Week, With a Price Tag

Module	What you built	This week	~2019
2	R&D-intensity panel, SEC FS	80 min lab	RA, ~1 week
3	EDGAR risk-factor pipeline + reusable skill	one morning	RA, 2–3 weeks
4	Filtered, verified Compustat funda pull	one afternoon	RA, several days
5	Hardened, rerunnable repo — paper trail, git	~3 hrs	RA, days of audit

The question this raises is not technical.

Place Your Dissertation on This Map




Pipeline stages $\times$ automatable $\leftrightarrow$ human-judgment $\times$ low $\leftrightarrow$ high value.

Do this now, silently

Place your own dissertation's stages on this map. Don't discuss yet — we open Block 1 with two volunteers defending their placement.

The Artifact and the Capital



Artifact: transferable, cheap to copy. Capital: built only by doing the work.

The personal audit

Can you discuss it without the LLM? Macnamara et al.: AI targets the sub-tasks where expertise gets built. Gerlich: attainment buffers decay, doesn't eliminate it.

Provenance: Who Is Accountable?

“A computer can never be held accountable.”

— IBM training manual, 1979

Nature's two rules (2023)

1. No LLM will be credited as an author — authorship carries accountability a model cannot discharge
2. Any use of these tools must be documented

A Retraction, Not a Citation

Retraction case study — labeled, never cited as evidence


A 2024 working paper claimed a randomized AI materials-discovery tool raised researcher output 44%

- Circulated to acclaim from prominent economists — then found fabricated
- Author severed from MIT; paper withdrawn
- The field's provenance-checking machinery failed *before* AI was writing papers

The open question

Does AI make that machinery stronger (audit trails) — or weaker (volume)?

Papers Are the Currency



Placement, tenure, prestige, salary all clear through publications. Bloom, Jones, Van Reenen & Webb (2020): ideas are getting harder to find. Lusher, Yang & Carrell (2023): more working papers the same week → worse outcomes for each, not just the weak ones.

Rules of Engagement

- Three blocks, ~30 min each — 5-min stretch between
- Each block: 2-min position statements (assigned Module 5) → open discussion → instructor synthesis
- No catastrophizing, no dismissing — every provocation paired with its counter
- Memo due one week after today

Block 1 — Knowledge vs. Learning

- If AI produces the artifact, what happens to the learning?
- Opens on your silent dissertation-placement — two volunteers defend it
- Which pipeline stages build intuition worth protecting?
- Accounting hook: AI-assisted audit judgment, the profession's parallel case

Closing question

Is PhD training a production process — or a transformation of the student?

Block 2 — Man vs. Machine

- Does it matter, epistemically or socially, who or what produced a result?
- Two identical extractions: RA-hand-coded vs. Claude-with-audit-trail — equivalent?
- Clean provenance can faithfully publish a wrong estimate
- 10-min detour: when a second model reviews the first model's code, what is actually verified?

Closing question

Does AI make the field's provenance-checking machinery stronger — or weaker?

Block 3 — Returns to Research

- Papers are the currency — what happens when production costs collapse?
- You built the Module 4 pipeline in an afternoon. Suppose every accounting PhD can.
- Scarce candidates: question taste, data access, credibility, accountability
- Homogenization: if every referee report sounds alike, what signal survives?

Closing question

Name one input you believe stays scarce — that's the bridge to the memo.

The Memo

“What I will still be paid for in 2036” — one page, due in one week

- Commit to specific components of the research production function that stay human
- Defend with one course reading *and* one experience from the build days
- Name one concrete investment for your remaining PhD years

Grading

Credit/no-credit on **defended specificity** — an optimist and a pessimist can both earn credit; a vague memo cannot.